

Responsible Machine Learning with Insurance Applications

Christian Lorentzen & Michael Mayer

Autumn 2022

Table of Contents

XAI: Introduction

XAI: Explaining Models

XAI: Improving Explainability

Introduction

How to explain and interpret a given model, even if it seems a black box?

Answering this question is a key aspect of responsible ML

1. Information for stakeholders
2. Detect problems in modeling process

XAI: e**X**plainable **A**rtificial **I**ntelligence

Collection of methods to explain and interpret models

Scope and Taxonomy

Scope

XAI methods for structured data and **bold** aspects below

Taxonomy of explainability

- ▶ **Global** vs. local: Describe model as a whole or around an observation.
- ▶ Model-specific vs. **model-agnostic**: Some methods are tailored to specific model classes (linear regression, tree-based), others work for all types of models.
- ▶ Intrinsic versus **post-hoc**: Simple models like a linear regression can be interpreted intrinsically, while complex models require post-hoc analysis of fitted model.

Notes

- ▶ Model-agnostic methods are always post-hoc
- ▶ Model-agnostic methods can also be applied to intrinsically interpretable models
- ▶ Won't make difference between “explainable”, “interpretable”, “intelligible”

XAI Outline

1. Introduction

- ▶ Notation
- ▶ Non-life insurance pricing
- ▶ Main example

2. Explaining Models

- ▶ Important post-hoc interpretation methods
- ▶ SHAP
- ▶ Improve GLM with the help of ML and XAI

3. Improving Explainability

Improve intrinsic explainability of complex models by simplifying their structure

Notation

Basic modeling situation

$$T(Y \mid \mathbf{x}) \approx m(\mathbf{x})$$

- ▶ Distributional property $T(Y \mid \mathbf{X} = \mathbf{x}) = T(Y \mid \mathbf{x})$ of response Y
- ▶ Model $m : \mathbf{x} \in \mathbb{R}^p \mapsto \mathbb{R}$ of feature vector \mathbf{X} with value $\mathbf{x} = (x^{(1)}, \dots, x^{(p)})^T$
- ▶ m estimated by \hat{m} from training data by minimizing objective criterion $\sum_{i=1}^n w_i L(\hat{y}_i, y_i) / \sum_{i=1}^n w_i + \lambda \Omega(m)$
- ▶ L : loss/scoring function, ideally strictly consistent for T ; $\lambda \Omega(m)$: optional penalty
- ▶ $\mathbf{w} = (w_1, \dots, w_n)^T$: vector of (optional) case weights
- ▶ $\mathbf{y} = (y_1, \dots, y_n)^T$: observed values of Y
- ▶ $\hat{\mathbf{y}} = (\hat{y}_1, \dots, \hat{y}_n)^T$: predicted/fitted values $\hat{y}_i = \hat{m}(\mathbf{x}_i)$
- ▶ $\mathbf{x}_1, \dots, \mathbf{x}_n$: n feature vectors; $x_i^{(j)}$: value of j -th feature of i -th observation

Examples of Models

- ▶ Linear regression
- ▶ Generalized linear models (GLM)
- ▶ Generalized additive models (GAM)
- ▶ Gradient boosted trees

Will peek into them as a quick refresher and to get used to notation

Linear Regression

- ▶ Model equation postulates

$$\mathbb{E}(Y \mid \mathbf{x}) = m(\mathbf{x}) = \beta_o + \beta_1 x^{(1)} + \dots + \beta_p x^{(p)}$$

- ▶ $(\beta_o, \beta_1, \dots, \beta_p)^T \in \mathbb{R}^{p+1}$: parameter vector to be estimated
- ▶ Objective: Minimize sum of squared errors

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

by linear least-squares

- ▶ Non-linear effects and interactions have to be added manually
- ▶ Penalized regression?
- ▶ Important extension: the generalized linear model (GLM)

Generalized Linear Model (GLM)

- ▶ Model equation postulates

$$\mathbb{E}(Y \mid \mathbf{x}) = m(\mathbf{x}) = g^{-1}(\eta(\mathbf{x})) = g^{-1}(\beta_o + \beta_1 x^{(1)} + \dots + \beta_p x^{(p)})$$

- ▶ g^{-1} : inverse link, g : link function, η : linear predictor
- ▶ Parameters β_j estimated by minimizing the (possibly weighted) *average deviance*

$$\bar{S}(\hat{m}, D_{\text{train}}) = \sum_{i=1}^n w_i S(\hat{y}_i, y_i) / \sum_{i=1}^n w_i$$

over training data $D_{\text{train}} = \{(y_i, w_i, \mathbf{x}_i), i = 1, \dots, n\}$

(Unit) deviance

- ▶ Distribution-specific measure: Poisson, Gamma, Bernoulli, normal, ...
- ▶ In our examples, we will often work with Poisson deviance
 $S(\hat{y}_i, y_i) = 2(y_i \log(y_i/\hat{y}_i) - (y_i - \hat{y}_i))$

Generalized Additive Model (GAM)

- ▶ Extension of the GLM
- ▶ Model equation assumes

$$\mathbb{E}(Y \mid \mathbf{x}) = m(\mathbf{x}) = g^{-1}(\beta_o + f_1(x^{(1)}) + \cdots + f_p(x^{(p)}))$$

- ▶ f_j : Sufficiently nice functions (some may be fully parameteric)
- ▶ Estimated to minimize average deviance, e.g. using backfitting
- ▶ Unlike a GLM, automatically accounts for non-linear effects
- ▶ Like a GLM, a GAM can also include interaction effects

Gradient Boosted Trees

- ▶ Typical black-box m
- ▶ Sum of decision trees
- ▶ In contrast to GAM, automatically picks up interactions
- ▶ Can optimize same objective criterion as GLMs/GAMs
- ▶ Using a different model structure and a different optimization technique
- ▶ Important implementations: LightGBM, XGBoost

Later, we will also work with neural nets

Non-Life Insurance Pricing

Main task: Predict *pure premium* of insurance policy

- ▶ Financial loss per year or per some other relevant exposure measure
- ▶ Used by company to optimize tariffs and to estimate expected future profit
- ▶ Predictions of statistical models fitted on historic data

Discussion

Why is it important to have good tariff?

Characterization of Insurance Policy

- ▶ $w > 0$: The exposure. Other quantities will refer to this
- ▶ N : Number of claims
- ▶ C : Total claim amount
- ▶ C/w : Pure premium
- ▶ $Y = N/w$: Claims frequency
- ▶ $Z = C/N$: Severity = avg cost per claim
- ▶ \mathbf{X} : One or more risk characteristics

Example (fictive motor third-part liability (MTPL) policies)

| id | w | N | C | C/w | Y | Z | Driver's age | Horse power |
|----|-----|-----|------|-------|-----|------|--------------|-------------|
| 1 | 1 | 0 | 0 | 0 | 0 | - | 28 | 80 |
| 2 | 0.5 | 2 | 5000 | 10000 | 4 | 2500 | 20 | 250 |
| 2 | 0.5 | 1 | 1000 | 2000 | 2 | 1000 | 21 | 250 |

Remark

Due to additivity of w , N , and C , these quantities can also be defined for multiple policies together, e.g., for the entire portfolio

Classic Pricing Models

- ▶ Instead of creating a model for $\mathbb{E}(C/w \mid \mathbf{x})$, decompose pure premium

$$C/w = (C/w) \cdot (N/N) = (N/w) \cdot (C/N) = YZ$$

into product of frequency Y and severity Z

- ▶ Frequency model: $\mathbb{E}(Y \mid \mathbf{x}) \approx m_Y(\mathbf{x})$
→ Poisson GLM with log link and case weights w
- ▶ Severity model: $\mathbb{E}(Z \mid \mathbf{x}) \approx m_Z(\mathbf{x})$
→ Gamma GLM with log link and case weights N , using only rows with $N > 0$
- ▶ Assuming conditional independence of Y and Z , pure premium model is then $\mathbb{E}(C/w \mid \mathbf{x}) \approx m_Y(\mathbf{x})m_Z(\mathbf{x})$

Alternative to GLMs

- ▶ Replace GLMs by GAMs or modern ML techniques
- ▶ Use same losses (deviance), weights, links

More on Non-Life Insurance Pricing

- ▶ The severity model can use different features than the frequency model
- ▶ The Gamma model with link is slightly biased → can be fixed by applying empirical multiplicative correction factor

$$c = \sum_{i=1}^n y_i / \sum_{i=1}^n \hat{y}_i$$

calculated on the training data

- ▶ Why using a log link for the Gamma model?
- ▶ As an alternative to model claims frequency using case weights w , directly model claim counts N without weights but using an offset of $\log(w)$
→ Same effects, different response, different evaluations

Main Example

Example (French motor third-part liability (MTPL) dataset)

1. Understand data
2. Descriptive analysis
3. Build claim frequency models

The models

- ▶ Poisson GLM
- ▶ Poisson GAM
- ▶ Poisson gradient boosted trees

Notes

- ▶ Grouped train/test split
- ▶ Model interpretation

Table of Contents

XAI: Introduction

XAI: Explaining Models

XAI: Improving Explainability

Introduction

Basic workflow to inspect and explain supervised learning model m

1. Study model performance
2. Study feature importance
3. Study feature effects, ideally also interactions

Focus

Global, model-agnostic, post-hoc explainability

Analysis result

- ▶ Information gain
- ▶ Reveal problems in model/data
- ▶ Increase confidence in model and modeler

Main references

- ▶ Online book of Christoph Molnar
- ▶ Tutorial (Mayer & Lorentzen 2020)

Chapter Outline

1. Software
2. Performance
3. Excursion: grouped data
4. Variable importance
5. Effects
6. Global surrogate models
7. Improve linear models by XAI
8. SHAP

Software for Post-Hoc Interpretation

R

- ▶ DALEX
- ▶ iml
- ▶ flashlight
- ▶ SHAP: kernelshap, shapviz, fastshap
- ▶ ...

Programming workflow

1. Build model
2. Create explainer object
3. Calculate and visualize results

Python

- ▶ scikit-learn inspect
- ▶ DALEX
- ▶ SHAP: shap
- ▶ ...

Example

Performance

- ▶ Study one or more relevant performance measures
- ▶ Often: Average loss or function of it
- ▶ Gives valuable information
- ▶ Training versus test performance? → assess overfitting/optimism
- ▶ Absolute and relative measures

Also helps to detect problems

Is performance **much lower** than expected?

- ▶ Preprocessing errors
- ▶ Missing key feature
- ▶ Convergence problem

Is it **much better**?

- ▶ Data partitions not independent?
- ▶ Leakage from response to feature?

Example: Claims Frequency Models

- ▶ Calculate weighted average Poisson deviance on test data:

$$\bar{S}(\hat{m}, D_{\text{test}}) = \sum_{i=1}^n w_i S(\hat{y}_i, y_i) / \sum_{i=1}^n w_i$$

with $S(\hat{y}_i, y_i) = 2(y_i \log(y_i/\hat{y}_i) - (y_i - \hat{y}_i))$

- ▶ Relative deviance improvement (one of many “pseudo R-squared”)

$$1 - \frac{\bar{S}(\hat{m}, D_{\text{test}})}{\bar{S}(\hat{m}_{\text{trivial}}, D_{\text{test}})},$$

where $\hat{m}_{\text{trivial}}(\mathbf{x}) = \sum_{D_{\text{train}}} w_i y_i / \sum_{D_{\text{train}}} w_i$ is the intercept-only model with constant predictions ideally calculated on the training data

- ▶ Repeat on training data (why?)

Excursion: Grouped Data

- ▶ Flawed validation strategy → biased performance assessment
- ▶ Detailed knowledge of data and model required → difficult to detect

Typical reason for flawed validation: Grouped data

- ▶ Pricing data
- ▶ Reserving: Models for ultimate claim amount
- ▶ Customer analytics: Browser behaviour of online visitors
- ▶ Banking: Financial transactions of clients

Grouped splitting

- ▶ Instead of random sampling of *rows*, we sample *groups*
- ▶ All rows of a group go in same data partition
- ▶ If ignored: Overfitting is being rewarded

Example and Simulation

Example (French MTPL)

- ▶ Model performance of GLM and boosted trees model without grouped splitting?
- ▶ Impact on model tuning?

Example (Simulation)

- ▶ Random data
- ▶ Linear regression and random forest
- ▶ 80%/20% split
- ▶ Increasing proportion of duplicated rows
- ▶ Random split versus grouped split

More on Grouped Data

- ▶ We used grouping structure to create clean data splits
- ▶ Sometimes, one wants to also make use of within-group info in model
- ▶ Panel data, time-series

Examples

- ▶ Insurance of large vehicle fleets: credibility factors
- ▶ Banking: Financial transactions of client

Tricky to have clean validation strategy and to apply model correctly

Variable Importance

1. Information: Most/least important features?
2. Challenge correctness of model
 - ▶ Results as expected or not?
 - ▶ Seemingly unimportant feature is top predictor → leakage?
 - ▶ Key features not among important features → preprocessing problem, not sufficient understanding of data or modeling situation?

Model-specific variable importance measures

- ▶ Linear model: normalized coefficients, test statistics etc.
- ▶ Tree-based models: Split gain or split count

Model-agnostic measures

- ▶ Permutation importance (Breiman 2001 for random forests)
- ▶ SHAP feature importance

Permutation Importance

Permutation importance of j -th feature $X^{(j)}$, data D , and performance measure \hat{S} :

$$\text{PVI}(j, D) = \hat{S}(\hat{m}, D^{(j)}) - \hat{S}(\hat{m}, D)$$

- ▶ $D^{(j)}$ is version of D with randomly permuted values in j -th feature column
- ▶ Read: How much \hat{S} drops after shuffling column j ? The higher, the more important. If 0, feature is unimportant

Algorithm to calculate $\text{PVI}(j, D)$ for all features

| Algorithm 1: Permutation importance |
|--|
| scoreOriginal \leftarrow performance on data for x <i>in</i> variables do <ul style="list-style-type: none">dataShuffled \leftarrow data with permuted column xscoreShuffled \leftarrow performance on dataShuffledimportance[x] \leftarrow scoreShuffled - scoreOriginal end output : importance |

Source: Mayer and Lorentzen, 2020

Remarks and Example

Remarks

- ▶ Computationally cheap → repeat m times
- ▶ Model is never refitted
- ▶ There is no formal definition of variable importance → inconsistency across methods
- ▶ Different definitions of permutation importance
- ▶ Strongly dependent features
→ decorrelate or analyze together
- ▶ Training or test data?

Example (French MTPL)

- ▶ PVI using exposure-weighted average Poisson deviance
- ▶ Hold-out data
- ▶ Compare with tree-split gain

Effects

Study and understand feature effects is of key importance

- ▶ How does $m(\mathbf{x})$ change with j -th feature?
- ▶ Often Ceteris Paribus: other components in \mathbf{x} fixed

Advantage of intrinsically interpretable models

- ▶ (Ceteris Paribus) effect of feature $X^{(j)}$ in a linear regression

$$\mathbb{E}(Y \mid \mathbf{x}) \approx m(\mathbf{x}) = \beta_o + \beta_1 x^{(1)} + \dots + \beta_p x^{(p)}$$

- ▶ In an additive model

$$\mathbb{E}(Y \mid \mathbf{x}) \approx m(\mathbf{x}) = \beta_o + f_1(x^{(1)}) + \dots + f_p(x^{(p)})$$

- ▶ In a black-box model?

Methods

1. Individual conditional expectation (ICE)
2. Partial dependence
3. Classic diagnostic plots
4. Interactions?
5. Later: SHAP dependence plots

Individual Conditional Expectation (ICE)

Basic thinking

- ▶ If m is additive in feature $X^{(j)}$, the Ceteris Paribus effect of $X^{(j)}$ is the same for all observations \rightarrow complete description of effect / full transparency
- ▶ If complex interactions involved \rightarrow approximate description only

Idea (Goldstein et al., 2015)

- ▶ Study Ceteris Paribus effect of $X^{(j)}$ for one observation
- ▶ *ICE function* for feature $X^{(j)}$ of model m and observation $\mathbf{x} \in \mathbb{R}^p$

$$\text{ICE}_j : v \in \mathbb{R} \mapsto m(v, \mathbf{x}_{\setminus j})$$

- ▶ $\mathbf{x}_{\setminus j}$ denotes all but the j -th component of \mathbf{x} , which is replaced by v
- ▶ *ICE curve* represents graph $(v, \text{ICE}_j(v))$ for grid of values $v \in \mathbb{R}$

Simple Algorithm

Algorithm to calculate $ICE_j(v)$

Algorithm 2: ICE for variable x and one observation

```
obs  $\leftarrow$  data row
for  $v$  in grid of values do
  | obs[ $x$ ]  $\leftarrow v$ 
  | ice[ $v$ ]  $\leftarrow$  prediction for obs
end
output : ice
```

Source: Mayer and Lorentzen, 2020

ICE Plot: Visualize ICE curves of multiple observations

Example

Notes

- ▶ Curves with different shapes indicate interaction effects
- ▶ Parallel curves \rightarrow additivity in $X^{(j)}$
- ▶ Centered ICE plots
- ▶ Usually on link scale (why?)
- ▶ ICE plots of higher dimension
- ▶ Training versus test data

Pros and Cons

- + Simple to compute
- + Easy to interpret (Ceteris Paribus)
- + Gives impression about interactions
- Suboptimal when Ceteris Paribus unnatural
- Model applied to rare/impossible \mathbf{x}
- Does not show what variables are interacting

Partial Dependence Plot PDP (Friedman 2001)

- ▶ Average of many ICE curves
- ▶ Ceteris Paribus effect of $X^{(j)}$ averaged over all interaction effects

Definition

- ▶ (Empirical) partial dependence function of feature j

$$\text{PD}_j(v) = \frac{1}{n} \sum_{i=1}^n \hat{m}(v, \mathbf{x}_{i,\setminus j})$$

- ▶ $\mathbf{x}_{i,\setminus j}$ feature vector of i -th observation without j -th component
- ▶ PDP equals graph $(v, \text{PD}_j(v))$ for grid of values $v \in \mathbb{R}$
- ▶ Sum runs over reference data (=?)

Algorithm

Calculate $PD_j(v)$ on grid of values for v

Algorithm 3: Partial dependence profile for variable x

```
 $n \leftarrow$  number of observations  
 $m \leftarrow$  grid size for variable  $x$   
iceProfiles  $\leftarrow$  matrix with  $n$  rows and  $m$  columns  
for  $i$  in 1 to  $n$  do  
  | iceProfiles[ $i$ ,:]  $\leftarrow$  ice curve for  $i$ th obs and variable  $x$   
end  
pd  $\leftarrow$  column means of iceProfiles  
output : pd
```

Source: Mayer and Lorentzen, 2020

Example

More on Partial Dependence

Remarks

- ▶ 2-dimensional PDP of $X^{(j)}$ and $X^{(k)}$:

$$PD_{jk}(v_j, v_k) = \frac{1}{n} \sum_{i=1}^n m(v_j, v_k, \mathbf{x}_{i, \setminus \{j, k\}})$$

- ▶ Accumulated local effects (ALE)

Pros and Cons

- + Simple to compute
- + Easy to interpret (Ceteris Paribus)
- Suboptimal when Ceteris Paribus unnatural
- Model applied to rare/impossible \mathbf{x}
- No information about interactions

Classic Diagnostic Plots

Related plots

1. Response versus covariate
→ Descriptive marginal effects
2. Predicted versus covariate
→ Modeled marginal effects
3. Residual versus covariate:
→ Bias assessment

Remarks

- ▶ Small and large datasets
- ▶ Binning of feature values
- ▶ Training versus test data?
- ▶ Relation to PDP?
- ▶ Pros and Cons?

Example

Interaction Effects

- ▶ Interactions: Linear models versus black-box models
- ▶ ICE plot for $X^{(j)}$ gives impression of total interaction effects associated with $X^{(j)}$

Pairwise interaction strength: Friedman's H

$$H_{jk}^2 = \frac{\sum_{i=1}^n \left[\text{PD}_{jk}(x_i^{(j)}, x_i^{(k)}) - \text{PD}_j(x_i^{(j)}) - \text{PD}_k(x_i^{(k)}) \right]^2}{\sum_{i=1}^n \left[\text{PD}_{jk}(x_i^{(j)}, x_i^{(k)}) \right]^2}$$

- ▶ Sums run over reference data
- ▶ Partial dependence functions are mean centered
- ▶ Interpretation of H^2 ? When close to 0 or 1?
- ▶ H versus H^2

Absolute Interaction Strength

Friedman's H is a relative measure \rightarrow absolute measure?

$$\tilde{H}_{jk} = \sqrt{\frac{1}{n} \sum_{i=1}^n \left[\text{PD}_{jk}(x_i^{(j)}, x_i^{(k)}) - \text{PD}_j(x_i^{(j)}) - \text{PD}_k(x_i^{(k)}) \right]^2}$$

Remarks on H and \tilde{H}

- ▶ Find out *how* features are interacting?
 \rightarrow Two-dimensional PDP, stratified PDP, SHAP
- ▶ Computational burden?
- ▶ Usually, one works on link scale

Example

Global Surrogate Models

Idea

- ▶ Fit intrinsically interpretable model m_I to predictions of \hat{m}
- ▶ Usually a small decision tree
- ▶ \hat{m}_I is called (global) surrogate model for \hat{m}
- ▶ Objective function and R-squared of \hat{m}_I ?

Remarks

- ▶ Training or test data?
- ▶ Variable importances of \hat{m}_I ?

Example

Improve Linear Models by XAI

Workflow

- ▶ Build strong GLM by the help of ML and XAI
- ▶ Why not directly use ML model?


Compare XAI aspects

1. Performance: Difference between GLM and ML model large or small?
2. Variable importance: Similar features important?
3. Main effects: Similar or not? → Change representation in GLM
4. Interaction effects: Add strong meaningful interactions to GLM

Example

SHAP: SHapley Additive exPlanations

- ▶ Local explanations
- ▶ Basic idea of SHAP?
- ▶ LIME



Shapley Values

Setting

- ▶ \mathcal{M} : Set of $p = |\mathcal{M}|$ players
- ▶ Playing cooperative game with numeric payoff
- ▶ Contribution of subset $\mathcal{L} \subseteq \mathcal{M}$ of players measured by function $v : \mathcal{L} \mapsto \mathbb{R}$

Question

How to distribute payoff fairly among the players?

Answer by Shapley (1953)

Player j should receive “Shapley value” = weighted average contribution

$$\phi_j(v) = \phi_j = \sum_{\mathcal{L} \subseteq \mathcal{M} \setminus \{j\}} \underbrace{\frac{|\mathcal{L}|!(p - |\mathcal{L}| - 1)!}{p!}}_{\text{Shapley weight}} \underbrace{(v(\mathcal{L} \cup \{j\}) - v(\mathcal{L}))}_{\text{Contribution of player } j}, \quad j = 1, \dots, p.$$

Fairness

ϕ_j averages the p average contributions of player j to coalitions of size $0 \leq |\mathcal{L}| \leq p-1$:

$$\text{Shapley weight} = \frac{|\mathcal{L}|!(p - |\mathcal{L}| - 1)!}{p!} = \frac{1}{p} \frac{1}{\binom{p-1}{|\mathcal{L}|}}.$$

Link to permutations?

Shapley values are only way to distribute total winnings fairly in the sense:

1. Efficiency: $v(\mathcal{M}) = \sum_{j=0}^p \phi_j$, where $\phi_o = v(\emptyset)$ denotes non-distributed payoff
2. Symmetry: If $v(\mathcal{L} \cup \{i\}) = v(\mathcal{L} \cup \{j\})$ for every $\mathcal{L} \subseteq \mathcal{M} \setminus \{i, j\}$, then $\phi_i = \phi_j$
3. Dummy player: If $v(\mathcal{L} \cup \{j\}) = v(\mathcal{L})$ for all coalitions $\mathcal{L} \subseteq \mathcal{M} \setminus \{j\}$, then $\phi_j = 0$
4. Linearity: Consider two cooperative games with gain functions v and w . Then, $\phi_j(v + w) = \phi_j(v) + \phi_j(w)$ and $\phi_j(\alpha v) = \alpha \phi_j(v)$ for all $1 \leq j \leq p$ and $\alpha \in \mathbb{R}$

Shapley Values in Statistics and ML

Early idea

Lipovetsky and Conklin (2001): Fair decomposition of *R-squared* in linear regression

Nowadays: Štrumbelj and Kononenko (2010, 2014), Lundberg and Lee (2017)

- ▶ Decompose *predictions* fairly into $m(\mathbf{x}) = \phi_o + \sum_{j=1}^p \phi_j$, $\phi_o = \mathbb{E}(m(\mathbf{X}))$
- ▶ Fair only if ϕ_j are Shapley values
- ▶ Natural contribution function: $v(\mathcal{L}) = m(\mathbf{x}_{\mathcal{L}})$, $\mathbf{x}_{\mathcal{L}}$ are components in $\mathcal{L} \subseteq \mathcal{M}$
- ▶ But: Features cannot be turned off in $m \rightarrow$ use statistics to estimate it

Situations where no estimation is required

- ▶ $p = 1$: Then $\phi_1 = m(\mathbf{x}) - \phi_o$
- ▶ Linear regression without correlations:

$$m(\mathbf{x}) = \underbrace{\beta_0}_{\phi_o} + \underbrace{\beta_1 x^{(1)}}_{\phi_1} + \dots + \underbrace{\beta_p x^{(p)}}_{\phi_p}$$

How to Estimate Contribution Function

Controversy in estimating $m(\mathbf{x}_{\mathcal{L}})$

- ▶ Statistically natural: Conditional expectation $\mathbb{E}(m(\mathbf{X}) \mid \mathbf{x}_{\mathcal{L}})$
- ▶ Causal inference prefers marginal expectations: $\mathbb{E}_{\mathbf{x}_{\mathcal{M} \setminus \mathcal{L}}}(m(\mathbf{X}))$

Algorithms (from slow to fast)

1. Monte Carlo sampling: For each j and many \mathcal{L} , contributions $v(\mathcal{L} \cup \{j\}) - v(\mathcal{L})$ are evaluated using marginal expectations and then plugged into Shapley's Eq.
2. Kernel SHAP: For many \mathcal{L} , evaluate $v(\mathcal{L})$ using marginal expectations. Then use weighted regression to get all Shapley values without plugging into Shapley's Eq.
3. TreeSHAP: Uses properties of trees to directly calculate $v(\mathcal{L})$ for all $\mathcal{L} \subseteq \mathcal{M}$ and then plugging into Shapley's Eq.

Example

From Local to Global Explanations

Notation

- ▶ X : $(n \times p)$ feature matrix with elements x_{ij} , $1 \leq i \leq n$, $1 \leq j \leq p$
- ▶ Φ : $(n \times p)$ matrix of SHAP values with elements ϕ_{ij}
- ▶ $\phi_o = \frac{1}{n} \sum_{i=1}^n \hat{m}(\mathbf{x}_i)$
- ▶ $\hat{m}(\mathbf{x}_i) = \phi_o + \sum_{j=1}^p \phi_{ij}$ for n feature vectors \mathbf{x}_i

Strategy to understand model as a whole

- ▶ SHAP feature importance: $I_j = \frac{1}{n} \sum_{i=1}^n |\phi_{ij}|$
- ▶ SHAP dependence plots: $\{(x_{ij}, \phi_{ij}), 1 \leq i \leq n\}$
- ▶ Interactions: Use x_{ik} , $k \neq j$ to add color to SHAP dependence plot (alternative?)

Examples

Each aspect separately and full analysis

SHAP Analysis to Improve Linear Model

Revisit our strategy

- ▶ A lot of info on a ML black-box can be generated very quickly
- ▶ Use it to build strong GLM

Example

Remember

SHAP has a solid theoretical foundation. In practice, some of it is lost because statistics is not mathematics.

Table of Contents

XAI: Introduction

XAI: Explaining Models

XAI: Improving Explainability

Introduction

“The best explanation of a simple model is the model itself” (Lundberg and Lee, 2017)

Two main strategies in XAI

- ▶ Model is black-box → interpret it post-hoc
- ▶ Make model less opaque / improve intrinsic explainability

Basic hierarchy of intrinsic explainability

1. Linear additive models (like GLMs)
2. Additive models (like GAMs)
3. Black-box models (like boosted trees or neural nets)

Note

- ▶ To link or not?
- ▶ Single decision tree

Boundaries are Blurred

Examples

- ▶ GLMs can have non-linear effects
- ▶ GLMs and GAMs can have interactions
- ▶ A complex GLM can be (almost) as black-box as a boosted trees model
- ▶ Boosted trees models can have all or some features additive
- ▶ Neural nets can have all or some features additive

Partly additive models = Additive with (possibly complex) interactions

- ▶ Additive time effects: $m(\mathbf{x}) = f(\text{Time}) + f'(\text{other features})$
- ▶ Additive gender effects: $m(\mathbf{x}) = f(\text{Gender}) + f'(\text{other features})$
- ▶ Additive model with non-additive location effects:
$$m(\mathbf{x}) = f_1(x^{(1)}) + \dots + f_{p_1}(x^{(p_1)}) + f(\text{location features})$$

Chapter Outline

Structuring boosted trees

- ▶ Additive models
- ▶ Partly additive models
- ▶ Monotonicity

Structuring neural nets

- ▶ Additive models
- ▶ Partly additive models

Tune boosted trees models

Only if time

Structuring Boosted Trees

Interpreting boosted trees models

- ▶ Single decision trees m_k are simple to interpret
- ▶ Boosted trees m are sums of K decision trees

$$m(\mathbf{x}) = m_1(\mathbf{x}) + \cdots + m_K(\mathbf{x})$$

- ▶ Interpretation of m only post-hoc

Will investigate two ways to structure boosted trees

1. Additive boosted trees
2. Partly additive boosted trees

Both are extremely useful in practice

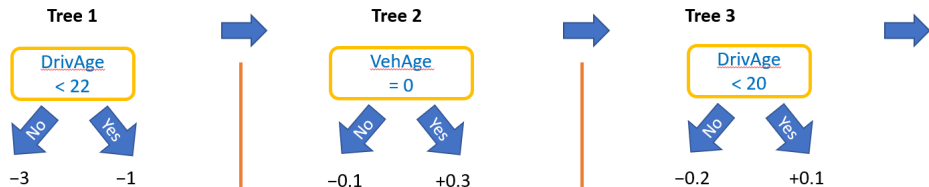
Additive Boosted Trees

What is a tree stump?

- ▶ Decision tree m_k with only one split (assume on j -th feature)
- ▶ Can be written as $m_k(\mathbf{x}) = v_1 + (v_2 - v_1)\mathbb{I}(x^{(j)} \leq s)$; ($v_1, v_2, s = ?$)

Boosted tree stumps are additive models

- ▶ $m(\mathbf{x}) = \beta_o + f_1(x^{(1)}) + \dots + f_p(x^{(p)})$
- ▶ $f_j, 1 \leq j \leq p$, are piecewise constant functions derived from m_k



More on Boosted Tree Stumps

- ▶ Additivity → full description of feature effects via ICE/PDP
- ▶ SHAP dependence plot? (Mayer 2022)
- ▶ Discussion: Pros and cons versus classic GAM?
- ▶ References: Lou et al. (2012), Nori et al. (2019)

Example

Partly Additive Boosted Trees

- ▶ Grow trees of depth $m = 2 \rightarrow$ pairwise interactions
- ▶ Partly additive model via *interaction constraints* (Lee et al., 2015)

Interaction constraints

- ▶ $IC = \{F_1, \dots, F_M\}$
- ▶ Each $F_m \subseteq \mathcal{M}$ is feature subset allowed to interact

How do they work?

- ▶ Consider a decision tree
- ▶ Rule: Each split considers features only from those F_m that contain all previous split variables of the branch.
- ▶ Thus, each branch will use features only from one $F_m \in IC$.
- ▶ Translates to tree and tree ensemble.

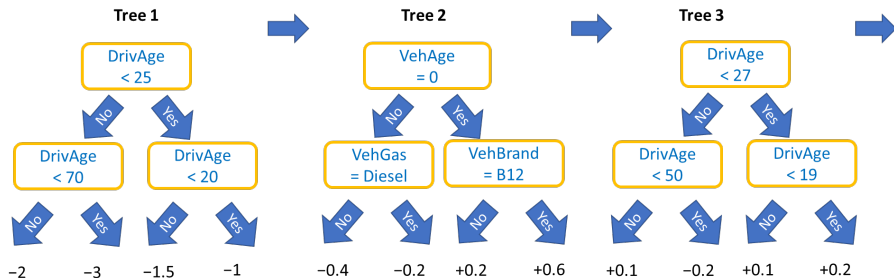
Partly Additive Models via Interaction Constraints

How to set IC so that model is additive in j -th feature?

- ▶ $F_m = \{X^{(j)}\}$ for some m
- ▶ $X^{(j)} \notin F_k$, for $k \neq m$

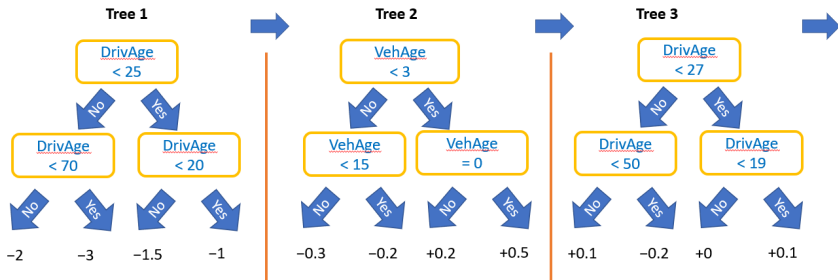
Example

$IC = \{\{\text{DrivAge}\}, \{\log\text{Density}\}, \{\text{PolicyRegion}\}, \{\text{VehAge}, \text{Brand}, \text{Gas}, \text{Power}\}\}$



More on Interaction Constraints

- ▶ If all elements in IC are disjoint, each tree uses features from only one F_m
- ▶ How is the first split variable determined?
- ▶ $IC = \{\{X^{(1)}\}, \dots, \{X^{(p)}\}\}$ gives additive model:



- ▶ Difference to boosted tree stumps?

Monotonic Constraints

- ▶ Monotonicity of $m(\mathbf{x})$ in j -th feature is another aspect of interpretability
- ▶ Violated natural monotonicity can have dramatic impact on trustworthiness
- ▶ Examples in car *collision* models?
- ▶ Simple to implement for decision trees!

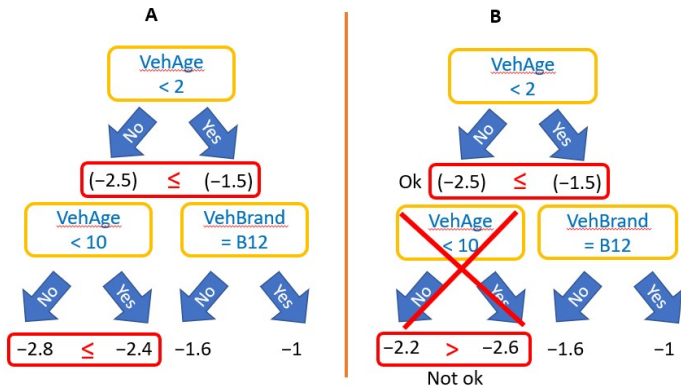


Figure: Aim: monotone decreasing predictions in vehicle age

More on Monotonic Constraints

- ▶ Monotonicity translates to tree ensembles
- ▶ Be careful when imposing monotonicity (why?)
- ▶ Can help to reduce wiggleness of effect
- ▶ Monotonicity for other model classes like GLMs, GAMs, neural nets?
- ▶ Monotonicity and outlying feature values

Example

Structuring Neural Nets

Swiss army knife of ML: Neural nets can

- ▶ mimic GLMs and GAMs,
- ▶ learn interactions and non-linear effects,
- ▶ fit data larger than RAM (e.g. images, videos),
- ▶ learn “online”,
- ▶ use multidimensional input and output,
- ▶ use input and output of mixed dimensionality,
- ▶ fit models with millions of parameters,
- ▶ perform non-linear dimension reduction,
- ▶ ...

How to create

1. linear,
 2. complex,
 3. additive, and
 4. partly additive
- neural nets?

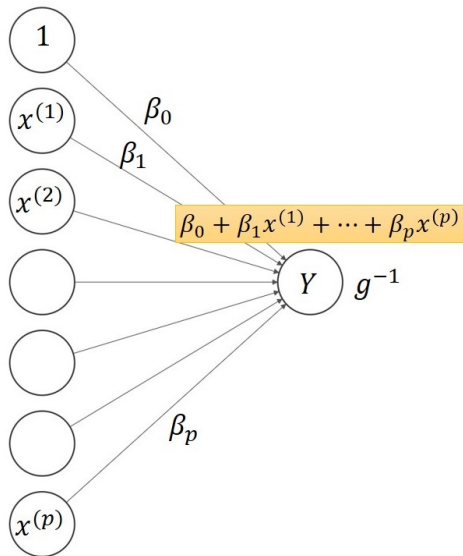
Some Notes on Neural Nets

- ▶ Why haven't we worked with neural nets so far?
- ▶ Neural nets versus boosted trees?
- ▶ References?
- ▶ TensorFlow, PyTorch, Keras
- ▶ Keras: sequential versus functional API
- ▶ Keras in R

A Simple Neural Net: GLM

Some slang

- ▶ Input and output layer?
- ▶ Nodes and node values?
- ▶ Fully connected / dense layer
- ▶ Exponential activation function



Example and Parameter Estimation

Parameters estimated by (mini-batch) gradient descent with backpropagation

1. Initialization: Randomly set parameters.
2. Forward step: Calculate predictions and their (total) loss of a *batch*
3. Backward step: Change parameters in the negative directions of the gradient of the loss to make loss smaller on that batch
4. Iterate Steps 2–3 for one epoch, i.e., until the full training data has been used
5. Iterate Step 4 until validation performance stops improving

Example

- ▶ Naive numeric representation of categoricals
- ▶ Callbacks?
- ▶ Feature scaling (why? how?)

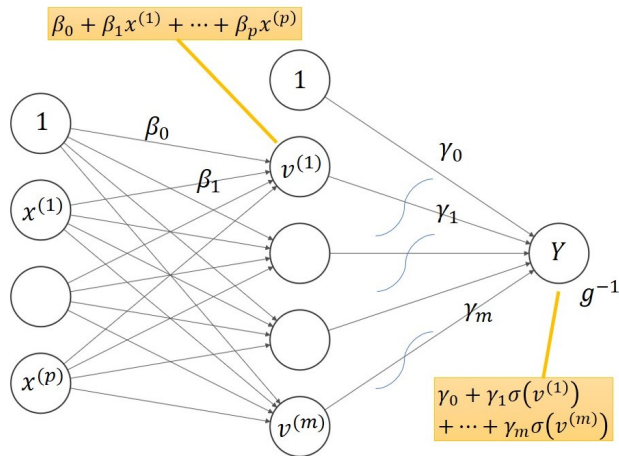
More Complex Models

Some additional slang

- ▶ Hidden layers
- ▶ Representational learning
- ▶ Activation functions: two purposes
- ▶ How to choose architecture?
- ▶ How to choose number of parameters/weights?

Example

- ▶ Three hidden layers
- ▶ 561 parameters

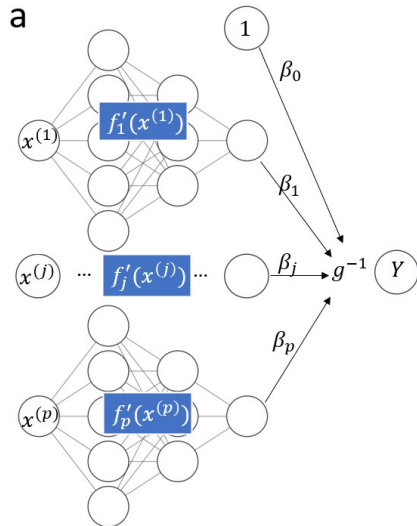


Additive Neural Nets (Agarwal et al. 2020)

- ▶ Represent each feature by single-output net
- ▶ Directly connected to output layer
- ▶ Linear components?
- ▶ Unordered categorical features?

Example

- ▶ 'VehBrand', 'PolicyRegion': 1-D embedding
- ▶ 'VehGas' and 'logDensity': Scaled and represented by linear function
- ▶ Rest: Scaled and represented by small net each
- ▶ Almost same structure as our original GAM



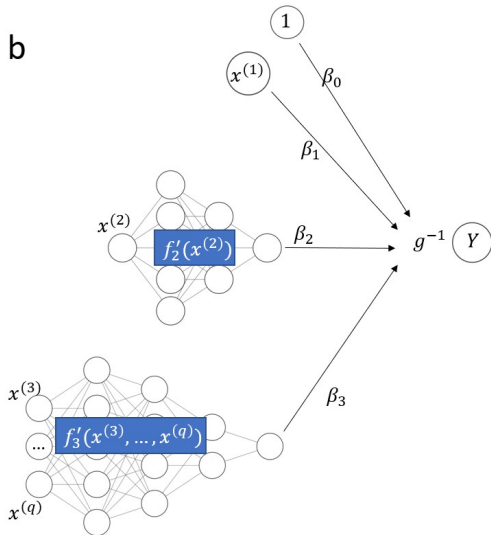
Source: Fig 1a in
<https://www.mdpi.com/1911-8074/15/5/193>

Partly Additive Models

- ▶ Pairwise interactions
- ▶ Partly additive model
- ▶ Example in geographic modeling
- ▶ Make model additive in driver features?

Example

- ▶ 'logDensity': Scaled and represented as a linear function
- ▶ 'DrivAge': Scaled and represented by a small sub-network
- ▶ 'PolicyRegion': 1-D embedding
- ▶ Vehicle features: sub-network with different inputs and five outputs



Source: Fig 1b in <https://www.mdpi.com/1911-8074/15/5/193>

Excursion: Tuning Boosted Trees

Model tuning in general

- ▶ How to choose hyperparameters of ML models?
- ▶ Each model class (GLMs, GAMs, random forests, boosted trees, neural nets, ...) has specialities that should be respected
- ▶ Examples?

Why focussing on boosted trees?

- ▶ Usually among best performing models for tabular data
- ▶ Boosting + SHAP → strong GLMs
- ▶ It needs some training

Aspects

1. Objective and evaluation metric
2. Number of boosting rounds
3. Learning rate
4. Further parameters

Objective and Metric

Ideal choice of loss function

- ▶ Meaningful for task
- ▶ Strictly consistent for target functional T

Translation to objective and metric

- ▶ Objective: average loss on training data (plus regularization) used for model training
- ▶ Evaluation metric: average (cross-)validation loss used for model comparison and selection

Example

Number of Boosting Rounds

Very important to select reasonable number of boosting rounds

- ▶ Boosting round = tree
- ▶ Too few rounds \rightarrow underfitting
- ▶ Too many rounds \rightarrow overfitting
- ▶ Heavily depends on choice of other parameters, thus difficult to choose

“Early stopping” as standard solution

- ▶ How does it work?
- ▶ Why is it so convenient?

Learning Rate

- ▶ Weight of each tree in final model
- ▶ Often between wide range of 1 and 0.005
- ▶ Good value heavily depends on number of boosting rounds
- ▶ Trick: select it so that early stopping ends after 100 – 1000 trees (why?)
- ▶ Halving the number of trees means doubling the learning rate for comparable performance

Regularization Parameters

Additional parameters to select

- ▶ number of leaf nodes
- ▶ tree depth
- ▶ loss penalties
- ▶ different types of subsampling rates
- ▶ ...

Choose them by (cross-)validation

- ▶ One by one
- ▶ Grid-search
- ▶ Random search

Note

- ▶ Early-stopping often compensates for suboptimal choice of other parameters
- ▶ Very different parameter combinations may lead to similar performance

Overall Strategy

Three steps

1. Choose strictly consistent and meaningful loss for functional T
→ objective and evaluation metric
2. Choose learning rate to get 100 – 1000 trees with early stopping
3. Select remaining parameters manually or by random search via (cross-)validation

Simplification

When to skip expensive Step 3?

Example

- ▶ French MTPL
- ▶ Speciality: grouped partitions